

Monte Carlo Algorithms

Dmytro Kedyk

Some Applications

- Quicksort and other randomized algorithms
- Random treap and other data structures
- Option pricing with binomial tree
- Quantifying and comparing system performance
- Fun interview questions

Learning Outcomes

- Understand the advantages of randomization
- Be able to generate random numbers and other objects
- Add simulation to your generic solution method toolbox
- Easily perform traditionally complex statistical analysis
- Correctly evaluate and compare system performance

Topics

- A quick review of probability
- Theory of randomization
- Generation of random objects
- Monte Carlo method

What is Probability

- Measures likelihood of **events** E_i
- E_i are subsets of a sample space S
 - E.g. S can be real numbers and E_i intervals
 - Samples $x \in S$ aren't events!
- Defined by axioms:
 - $\text{Prob}(E_i) \geq 0$
 - $\text{Prob}(S) = 1$
 - For disjoint events $\text{Prob}(\cup E_i) = \sum E_i$

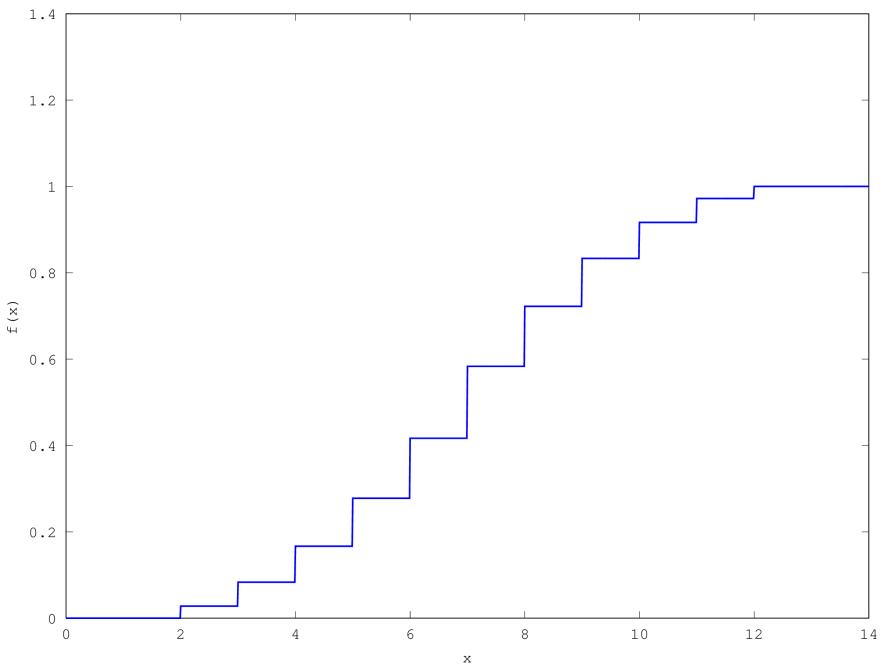
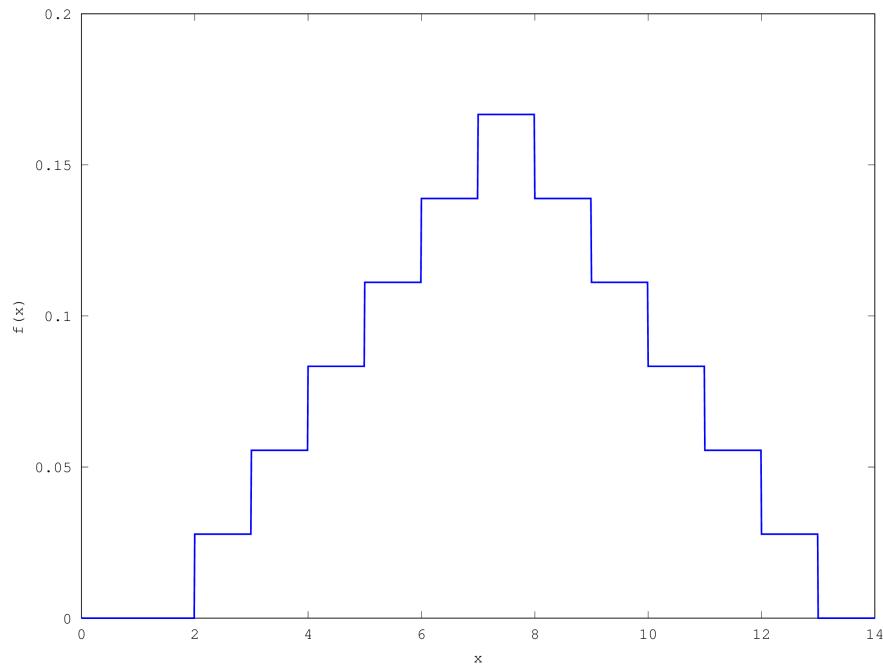


What is a Distribution

- Uses **probability density function (PDF)** $f(x)$ to assign values to $x \in S$
 - $\forall E \text{ Prob}(E) = \int_{x \in E} f(x)$
 - Values $f(x)$ need not be probabilities!
- For connected range events **cumulative distribution function (CDF)** $F(x) = \int_{\infty < t \leq x} f(t)$ directly gives probabilities

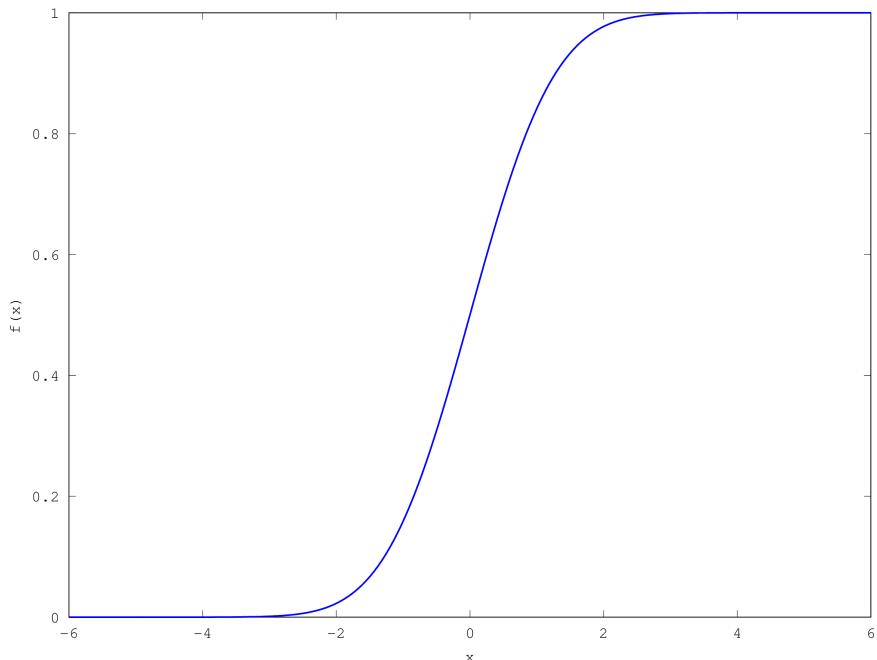
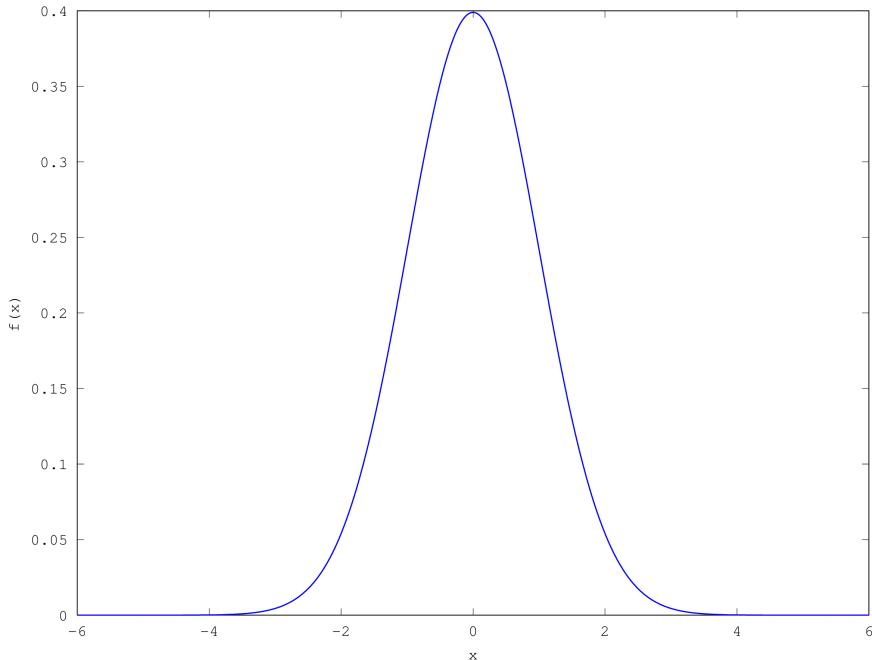
Die Sum Distribution

- $f(x) = \max(0, (6 - |7 - x|)/36)$
- $F(8) - F(5) = \text{Prob}(\text{sum is } 6, 7, \text{ or } 8)$



Normal(μ , σ) Distribution

- μ is mean, σ standard deviation
- Models many events, heavily used in statistics
- $f(x) = \exp(-((x - \mu)/\sigma)^2/2)/(\sigma\sqrt{2\pi})$

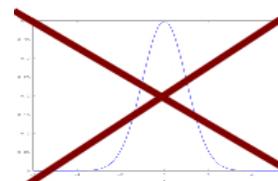


Evaluating Normal CDF

- E.g. two-sided confidence interval is $2F(x) - 1$
- $F(x) = 0.5 + \operatorname{erf}(x/\sqrt{2})/2$
- Can't express erf in terms of elementary functions, must approximate:
 - $\operatorname{erf}(x) \approx 1 - (1 + \sum_{0 \leq i < 6} a_i x^i)^{-16}$
 - a_i are $0.0705230784, 0.0422820123,$
 $0.0092705272, 0.0001520143, 0.0002765672,$
 0.0000430638
 - approximation has $\leq 10^{-7}$ error

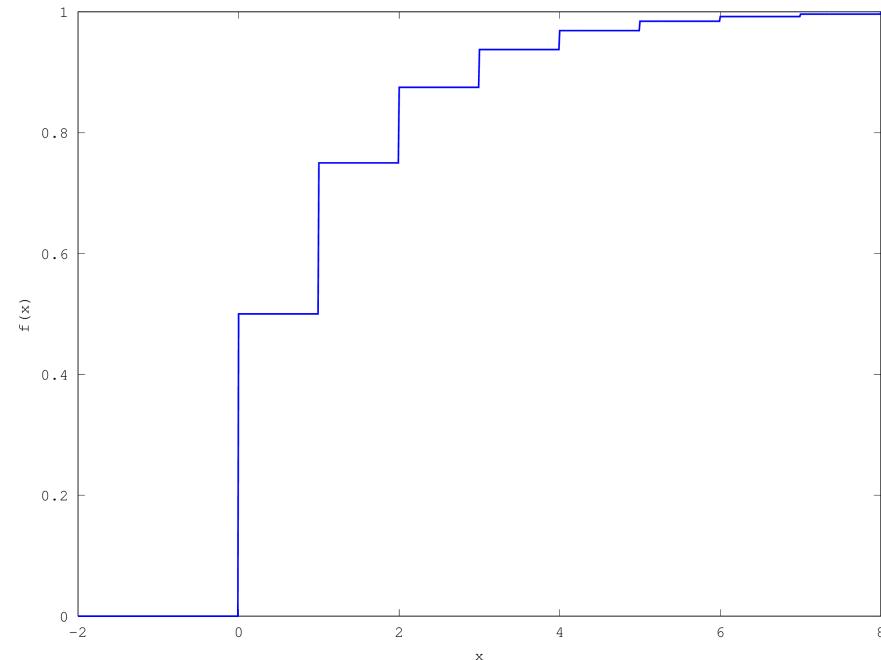
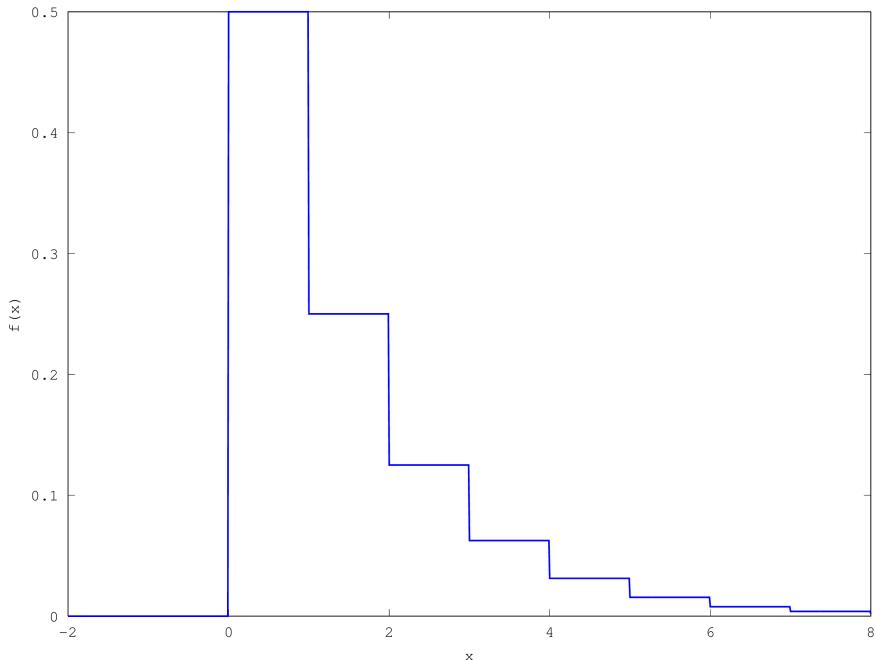
Using Distributions

- Often don't know the right distribution!
 - Every distribution is a good or bad model for some events
 - E.g. normal is a bad model for flying bird's avoiding a pole
- In many cases need only a summary
 - **Expected value** is $\bar{x} = \int_{x \in S} xf(x)$
 - **Variance** is $E[(x - \bar{x})^2]$
- Summaries may lose information
 - E.g. for multi-peak data μ doesn't mean much



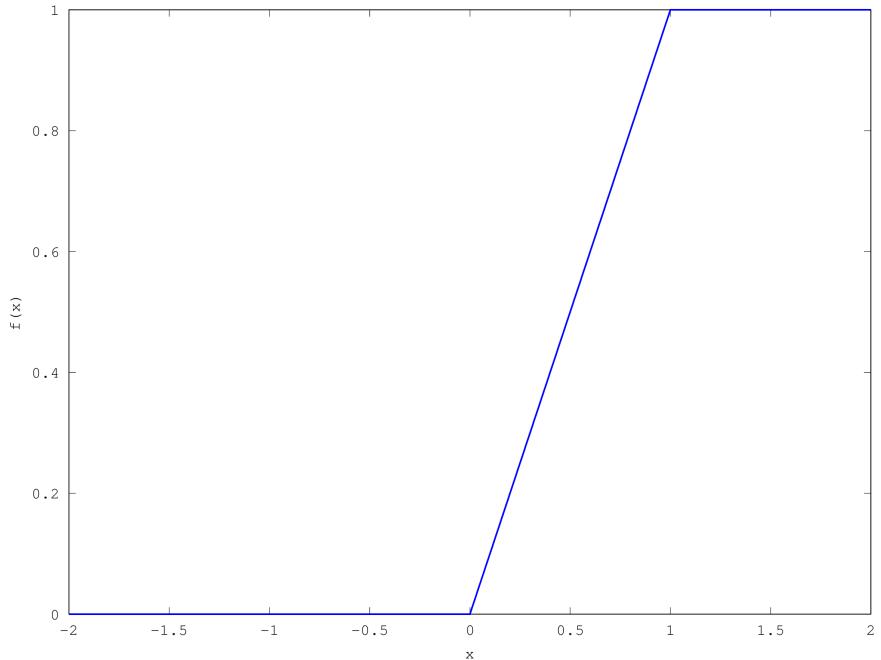
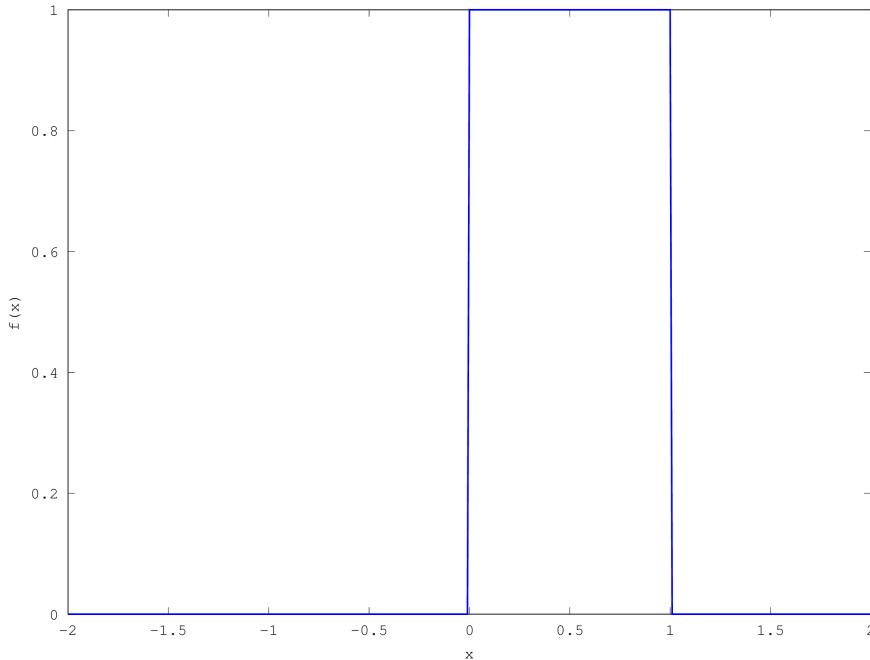
Geometric(p) Distribution

- How long it takes for p -coin to give “tails”
- x discrete, $f(x) = (1 - p)^{x-1}p$
- Models waiting for specific event, $E[x] = 1/p$



Uniform(a, b) Distribution

- Easy to sample from, basic building block
- $f(x) = 1/(b - a)$ if $a \leq x \leq b$ and 0 otherwise
- Used for no information models



Topics

- A quick review of probability
- Theory of randomization
- Generation of random objects
- Monte Carlo method

Randomization Helps

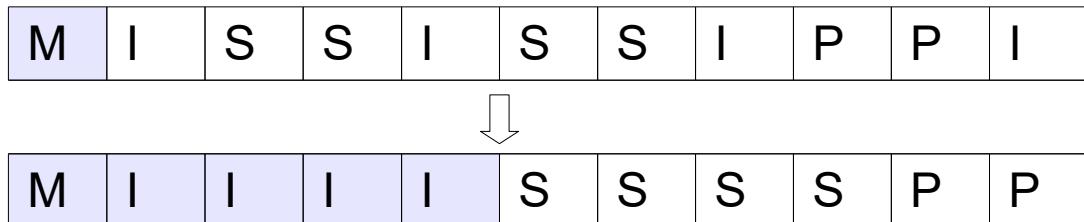
- In every game there is an optimal **randomized** strategy
- But may not be optimal **deterministic** one
- E.g. rock-paper-scissors—random choices draw on average
- But can learn to defeat a clever deterministic strategy, e.g. using machine learning

Randomized Algorithms

- Game between algorithm and input
 - Using `rand()` may avoid bad worst cases on average
- **Randomness in data \neq rand()**
- But may help too
 - The worst-case input may be unlikely
 - So many algorithms are fast in practice and slow in theory
 - E.g. unbalanced binary tree

Quicksort

- Pick a pivot, split array into $<$ pivot and \geq pivot, and recurse on each half



- How to pick?—first, last, median of 3 may give $O(n^2)$ runtime
- Random!
 - $O(n \ln(n))$ runtime on average
 - Also **tail inequality**— $\text{Prob}(\text{runtime} > O(n \ln(n)))$ is exponentially small!

Types of Randomized Algorithms

- **Las Vegas**—expected performance
 - E.g. random pivot quicksort— $E[\text{runtime}] = O(n \ln(n))$
- **Monte Carlo**—expected correctness
 - E.g. Miller-Rabin primality test—tiny chance of error
 - Repeat to reduce error
- Often randomized algorithms are the fastest known

Topics

- A quick review of probability
- Theory of randomization
- Generation of random objects
- Monte Carlo method

What is a Random Number Sequence?

- Hard to define exactly, but:
- A sequence is random if can't predict next value with probability > that of a guess
- Is a sequence containing 10^9 consecutive 0's random?

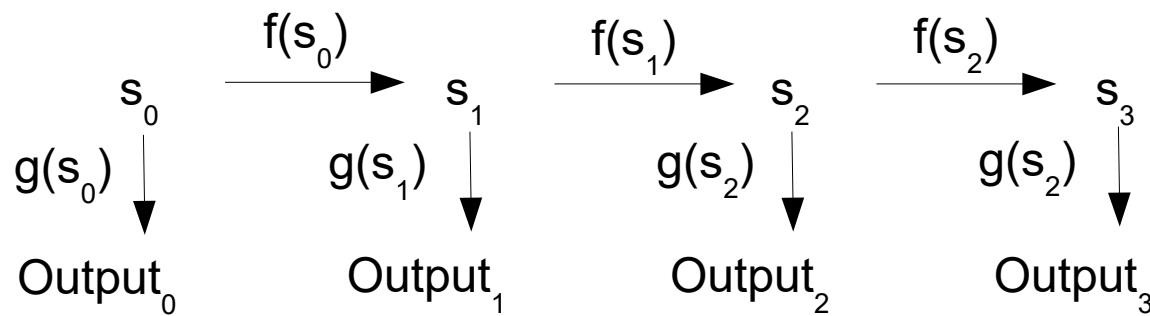


Getting Random Numbers

- Hire a coin flipper
- Measure physical phenomena, e.g., atmospheric noise or radioactive decay
- Record CPU or hard drive activity, mouse movement, keyboard actions, etc.
- Slow and unportable!

Pseudorandom Numbers

- Generated deterministically and indistinguishable from random
- General algorithm:
 - Start from some initial state
 - Output and the next state are functions of the current state



Generator Quality

- Produce random-looking outputs
 - **Period** $\geq 2^{64}$ —after how long the sequence repeats itself
 - High **equidistribution**—largest k for which consecutive k values can be any values
 - Pass statistical tests—can't reject hypothesis that the sequence is random

Generator Requirements

- Return random *double* $u \in (0, 1)$, not $[0, 1]$, to make $\log(u)$ well-defined
- Simple, fast, and portable
- Optionally generate independent streams for parallelization

Linear Congruential Generator

- Single-word state s
- Transition $s = (as + c) \% m$
 - a, c, m are picked constants
- $u = (s + 1)/(m + 1)$
- E.g. $a = 2685821657736338717$, $m = 2^{64}$, $c = 0$,
 $s_0 = 123456789$
 - $s_1 = 8624929095735532502$
 - $u_1 = 0.46755834315649508$

Linear Congruential Generator

- Don't use as is:
 - Period of the lower k bits = 2^k
 - Fails many tests
 - Multiplication needs double precision to avoid overflow

Picking the Initial State

- Function of system time and a password – fast, simple, portable, and very random
- Operating system random source—for cryptography
- Restored generator state from the last run—complete independence of generator runs

Xorshift

- Much better generator, transition
 - Interprets state as Boolean vector
 - And multiplies it by a Boolean matrix
- The matrix has a special form—can implement using shift and xor

Xorshift Code

```
class Xorshift
{
    unsigned int state;
    enum{PASSWORD = 19870804};

public:
    Xorshift(unsigned int seed = time(0) ^ PASSWORD)
    {
        assert(numeric_limits<unsigned int>::digits == 32);
        state = seed ? seed : PASSWORD;
    }

    static unsigned int transform(unsigned int x)
    {
        x ^= x << 13;
        x ^= x >> 17;
        x ^= x << 5;
        return x;
    }

    unsigned int next() {return state = transform(state);}
    double uniform01() {return 2.32830643653869629E-10 * next();}
};
```

Xorshift Properties

- 13, 17, and 5 are picked by theory, exhaustive search, and testing
- Minor quality problems
 - The period is $2^{32} - 1$ (state is never 0)—too small
 - Matrix multiplication is a linear operation—bits of successive numbers are correlated
- To remove correlation and increase period
 - Use 64-bit state, this changes constants
 - Combine with a simple LCG

Improved Xorshift

```
class QualityXorshift64
{
    unsigned long long state;
    enum{PASSWORD = 19870804};

public:
    QualityXorshift64 (unsigned long long seed =
        time(0) ^ PASSWORD)
    {
        assert(numeric_limits<unsigned long long>::digits == 64);
        state = seed ? seed : PASSWORD;
    }
    static unsigned long long transform(unsigned long long x)
    {
        x ^= x << 21;
        x ^= x >> 35;
        x ^= x << 4;
        return x * 2685821657736338717ull;
    }
    unsigned long long next() {return state = transform(state);}
    double uniform01() {return 5.42101086242752217E-20 * next();}
};
```

Improved Xorshift Properties

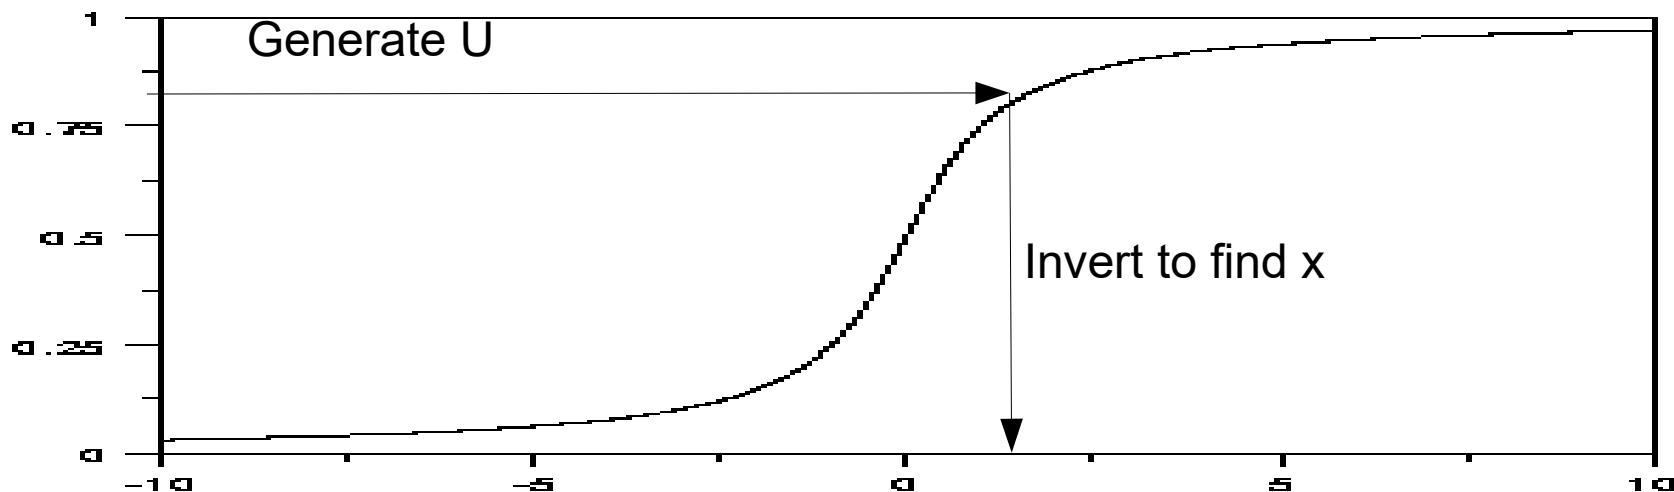
- Doesn't generate 0, so $u \in (0, 1)$
- Very fast, passes most tests
- Period 2^{64} is long enough for all practical uses
- Use as default generator unless have another good one from some API
- Can use the transition function for hashing

Other Good Generators

- MRG32k3a
 - Slower, but passes all tests and supports independent streams
- RC4
 - Much slower, but cryptographically secure
- Mersenne Twister
 - Same speed and test performance, and huge state
 - But complex implementation and uses 4KB memory

Samples from Probability Distributions

- **Inverse method**—cumulative distribution function F is a function $x \rightarrow [0,1]$, so $F^{-1}(u)$ is a random variate
- Can calculate F^{-1} numerically



Inverse Method Example

- Exponential distribution with parameter λ : $F(x) = 1 - e^{-\lambda x}$
 - Solve $F(x) = u$ to get $x = -\ln(1 - u)/\lambda$
 - Simplifies to $-\ln(u)/\lambda$ because u and $1 - u$ have the same distribution
- Many other methods for generation, most are distribution-specific
 - Boost has a good API

Some Continuous Generators

- Uniform(a, b) with $a < b$:

```
double uniform(double a, double b)
{return a + (b - a) * uniform01();}
```

- Exponential – memoryless waiting times:

```
double exponential01() {return
-log(uniform01());}
```

- Many more complex ones
 - Normal, gamma, Cauchy, etc.

Some Discrete Generators

- Bernoulli(p) – 1 with prob p and 0 with $1 - p$:

```
bool bernoulli (double p) { return  
uniform01 () <= p; }
```

- Geometric(p)—the number of times Bernoulli(p) = 0 before it's 1:

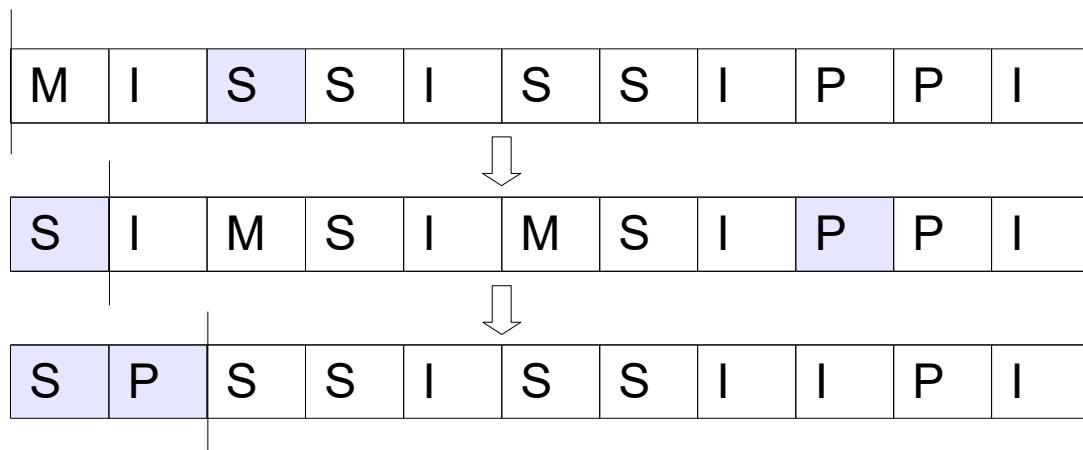
```
int geometric (double p)  
{  
    int result = 0;  
    while (!bernoulli (p)) ++result;  
    return result;  
}
```

Example Output

```
GlobalRNG.uniform01() 0.904229
GlobalRNG.uniform(10, 20) 10.2389
GlobalRNG.normal01() 0.508248
GlobalRNG.normal(10, 20) -8.47315
GlobalRNG.exponential01() 1.42522
GlobalRNG.gamma1(0.5) 0.00646845
GlobalRNG.gamma1(1.5) 0.355508
GlobalRNG.weibull1(20) 0.894177
GlobalRNG.erlang(10, 2) 13.6015
GlobalRNG.chisquared(10) 10.7126
GlobalRNG.t(10) 0.492003
GlobalRNG.logNormal(10, 20) 5.51916e+010
GlobalRNG.beta(0.5, 0.5) 0.205863
GlobalRNG.F(10, 20) 2.49824
GlobalRNG.cauchy01() 0.585426
GlobalRNG.binomial(0.7, 20) 14
GlobalRNG.geometric(0.7) 0
GlobalRNG.poisson(0.7) 2
```

Generating Random Objects

- Permutations—swap the first element with a random one and recursively permute the remaining $n - 1$
- Combinations of k out of n —use above for k steps



Topics

- A quick review of probability
- Theory of randomization
- Generation of random objects
- Monte Carlo method

Law of Large Numbers

- Average of many enough samples is the true average
 - Given n iid samples X_i from distribution T such that $E[X_i] = M$, $(\sum X_i)/n = \bar{x} \rightarrow M$ for $n \rightarrow \infty$
 - $\bar{x} \rightarrow$ sample mean of T
 - $s = (\sum(X_i - \bar{x})^2)/(n - 1) \rightarrow$ sample variance of T
 - Divisor for s is $n - 1$
 - Because variance = $E[X_j - (\sum_{i \neq j} X_i)/n]^2$

Central Limit Theorem

- Essentially LLN with error bounds
 - Given n iid samples X_i from a distribution with mean μ and finite variance σ^2 , for $n \rightarrow \infty$, $\bar{X} \sim \text{normal}(\mu, \sigma^2/n)$
 - LLN and technical Slutsky's theorem allow using s instead of σ
 - μ is $\bar{X} \pm 2\sqrt{s/n}$ with 95% probability!
 - The multiplier is 1.96 actually, but use 2 for convenience, or others

Monte Carlo Idea

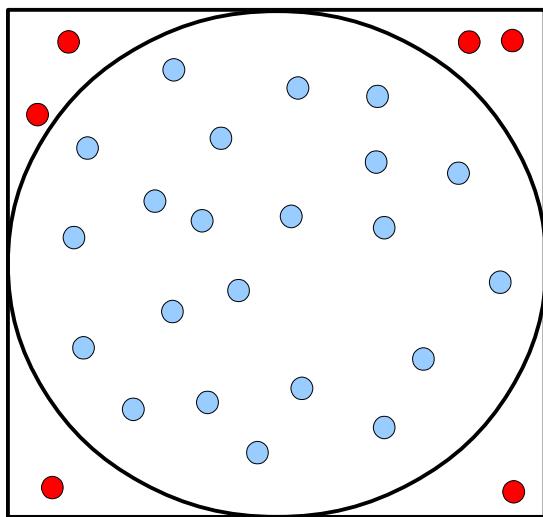
- Define quantity of interest X and compute it using the CLT
 - Need function f producing iid events with value X_i , such that $E[X_i] = X$
- Important events must be generated often enough
 - Otherwise need very large n for the CLT to kick in
 - Rare event problem!

Monte Carlo Algorithm

- Define X , pick f and a confidence level
- Until out of patience or (n is large enough and error small enough)
 - $X_i \leftarrow f()$
 - Incrementally update \bar{x} and s with X_i
- Return $X \leftarrow \bar{x} \pm$ the error

Example—Computing π

- Area of a circle with radius r is πr^2 and of its enclosing square $4r^2$
- $X = \pi/4 = (\text{the area of the circle})/(\text{the area of the square})$



$$\pi \approx 4 \times 22/(22 + 6) \approx 3.142$$

Computing π

- f generates random points $p \in (-1, 1) \times (-1, 1)$ and returns $X_i = (\text{distance}(p, (0, 0)) \leq 1)$
- $X_i = 1$ and $X_i = 0$ happen often enough
- After 10^8 2D uniform($-1, 1$) variates, $\pi = 3.14182 \pm 0.000493$ with 99.73% confidence

Calculating Mean and Error Incrementally

- Store and update $\sum X_i$ and $\sum X_i^2$
- After $n > 1$ simulations
 - $\bar{x} = \sum X_i / n$
 - $s = \max((\sum X_i^2 - \bar{x}^2 / n) / (n - 1), 0)$
 - Need *max* for numerical issues!
 - Variance of mean = s/n

Monte Carlo Good and Bad

- $O(n)$ time and $O(1)$ space
- A simulated event can produce k values
 - Effectively perform k related simulations at cost of one
- $O(1/\sqrt{n})$ convergence is too slow
 - Variance reduction via **common random numbers**
 - fix everything that isn't simulated
 - E.g. when simulating performance of a randomized algorithm, run all simulations on the same input

Predict World Cup Winner

- Highest rated team may not be most likely winner
 - Team with relatively easiest opponents is
- Team ratings R_i determine game result probabilities
 - $\text{Prob}(A \text{ wins against } B) = 1/(1 + 10^D)$, where $D = (R_B - R_A)/400$
 - Use this to simulate the tournament tree
- Expected winner is the most frequent winner

References

- Kedyk, D. (2020). *Implementing Useful Algorithms and Data Structures in C++: Toward Production Code*. Independently Published.

Lab

- Estimate π correctly to 3 decimal places
 - Use any programming language
 - That has a simple random number generator
 - Beware: accuracy of $\pi \neq$ accuracy of $\pi/4$
 - Multiplication by 4 magnifies the error